

# Breast Cancer Detection Kaidoko round 2 AI Developer Intern

by Ishan Sharma

we are firstly going to import all the necessary libraries at first. We may also need additional libraries so we will add those during that period only

```
In [1]: ▶ import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

sns.set()
plt.style.use('ggplot')
```

```
In [2]: ▶ df = pd.read_csv("C:\\Users\\ishan\\Downloads\\breast-cancer.csv")
df.head()
```

Out[2]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

5 rows × 32 columns

```
In [3]: ▶ df.diagnosis.unique()
```

Out[3]: array(['M', 'B'], dtype=object)

```
In [4]: ▶ # M is Malignant or cancerous
# B is Benign or non cancerous
```

## DATA PREPROCESSING

in this we usually have a set of processes which includes describe, it helps in looking the overview of the data points and help in looking at the outliers or help in suspicion (11.7 in radius mean are not very common so anything significantly less could be outlier or help us)

In [5]: `df.describe()`

Out[5]:

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothnes
<b>count</b>	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569
<b>mean</b>	3.037183e+07	14.127292	19.289649	91.969033	654.889104	C
<b>std</b>	1.250206e+08	3.524049	4.301036	24.298981	351.914129	C
<b>min</b>	8.670000e+03	6.981000	9.710000	43.790000	143.500000	C
<b>25%</b>	8.692180e+05	11.700000	16.170000	75.170000	420.300000	C
<b>50%</b>	9.060240e+05	13.370000	18.840000	86.240000	551.100000	C
<b>75%</b>	8.813129e+06	15.780000	21.800000	104.100000	782.700000	C
<b>max</b>	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	C

8 rows × 31 columns

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    569 non-null    int64
1   diagnosis                            569 non-null    object
2   radius_mean                          569 non-null    float64
3   texture_mean                         569 non-null    float64
4   perimeter_mean                       569 non-null    float64
5   area_mean                           569 non-null    float64
6   smoothness_mean                      569 non-null    float64
7   compactness_mean                     569 non-null    float64
8   concavity_mean                       569 non-null    float64
9   concave points_mean                  569 non-null    float64
10  symmetry_mean                        569 non-null    float64
11  fractal_dimension_mean               569 non-null    float64
12  radius_se                            569 non-null    float64
13  texture_se                           569 non-null    float64
14  ...
```

In [7]: `df.corr()`

Out[7]:

	id	radius_mean	texture_mean	perimeter_mean	area_r
id	1.000000	0.074626	0.099770	0.073159	0.09
radius_mean	0.074626	1.000000	0.323782	0.997855	0.98
texture_mean	0.099770	0.323782	1.000000	0.329533	0.32
perimeter_mean	0.073159	0.997855	0.329533	1.000000	0.98
area_mean	0.096893	0.987357	0.321086	0.986507	1.00
smoothness_mean	-0.012968	0.170581	-0.023389	0.207278	0.17
compactness_mean	0.000096	0.506124	0.236702	0.556936	0.49
concavity_mean	0.050080	0.676764	0.302418	0.716136	0.68
concave points_mean	0.044158	0.822529	0.293464	0.850977	0.82
symmetry_mean	-0.022114	0.147741	0.071401	0.183027	0.15

In [8]: `df.drop('id', axis=1, inplace=True)`

In [9]: `df.isnull().sum()`

Out[9]:


diagnosis	0
radius_mean	0
texture_mean	0
perimeter_mean	0
area_mean	0
smoothness_mean	0
compactness_mean	0
concavity_mean	0
concave points_mean	0
symmetry_mean	0
fractal_dimension_mean	0
radius_se	0
texture_se	0
perimeter_se	0
area_se	0
smoothness_se	0
compactness_se	0
concavity_se	0
concave points_se	0
symmetry_se	0
fractal_dimension_se	0

no missing value in this which is a good thing

## encoding categorical data

```
In [10]: df['diagnosis'] = df['diagnosis'].apply(lambda val:1 if val=='M' else 0)
#we did [1,0] as it numericals are benefiary in model training and testing
```

```
In [11]: plt.figure(figsize=(20,12))
corr=df.corr()
mask = np.triu(np.ones_like(corr, dtype=bool))
sns.heatmap(corr, mask=mask, linewidths=1, annot=True, fmt = ".2f")
plt.show()
```



```
In [12]: corr_matrix = df.corr().abs()
mask = np.triu(np.ones_like(corr_matrix, dtype=bool))
tri_df = corr_matrix.mask(mask)

to_drop = [x for x in tri_df.columns if any(tri_df[x]>0.92)]

df = df.drop(to_drop, axis=1)

print(df.shape)
```

(569, 23)

In [13]: `df.head()`

Out[13]:

	diagnosis	texture_mean	smoothness_mean	compactness_mean	concave points_mean	symmetr
0	1	10.38	0.11840	0.27760	0.14710	
1	1	17.77	0.08474	0.07864	0.07017	
2	1	21.25	0.10960	0.15990	0.12790	
3	1	20.38	0.14250	0.28390	0.10520	
4	1	14.34	0.10030	0.13280	0.10430	

5 rows × 23 columns

## Buliding Model

In [14]: `X=df.drop('diagnosis', axis=1)  
y=df['diagnosis']`

In [15]: `from sklearn.model_selection import train_test_split  
X_train, X_test, y_train ,y_test =train_test_split(X,y, test_size=0.2, ran`

In [16]: `# scaling data  
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)`

In [17]: `X_train.shape`

Out[17]: (455, 22)

## LogisticRegression

```
In [18]: from sklearn.linear_model import LogisticRegression
log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
```

```
Out[18]: LogisticRegression
LogisticRegression()
```

```
In [19]: y_pred = log_reg.predict(X_test)
```

```
In [20]: y_pred
```

```
Out[20]: array([1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1,
0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0,
0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0,
1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0,
1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
0, 1, 1, 0], dtype=int64)
```

```
In [21]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
print(accuracy_score(y_train, log_reg.predict(X_train)))
log_reg_acc = accuracy_score(y_test, log_reg.predict(X_test))
print(log_reg_acc)
y_pred = log_reg.predict(X_test)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0.989010989010989
```

```
0.9649122807017544
```

```
[[66  1]
 [ 3 44]]
```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	67
1	0.98	0.94	0.96	47
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

## Decision Tree

```
In [22]: from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV

dtc = DecisionTreeClassifier()

parameters = {
    'criterion': ['gini', 'entropy'],
    'max_depth': range(2, 32, 1),
    'min_samples_leaf': range(1, 10, 1),
    'min_samples_split': range(2, 10, 1),
    'splitter': ['best', 'random']
}

grid_search_dt = GridSearchCV(dtc, parameters, cv=5, n_jobs=-1, verbose=1)
grid_search_dt.fit(X_train, y_train)
```

Fitting 5 folds for each of 8640 candidates, totalling 43200 fits

```
Out[22]:
GridSearchCV
  estimator: DecisionTreeClassifier
    DecisionTreeClassifier
```

```
In [23]: grid_search_dt.best_params_
```

```
Out[23]: {'criterion': 'gini',
'max_depth': 6,
'min_samples_leaf': 1,
'min_samples_split': 7,
'splitter': 'random'}
```

```
In [24]: grid_search_dt.best_score_
```

```
Out[24]: 0.9626373626373625
```

```
In [25]: dtc = DecisionTreeClassifier(criterion='entropy', max_depth=15, min_sample
```

```
In [26]: dtc.fit(X_train, y_train)
```

```
Out[26]:
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=15, min_samples_l
eaf=4,
min_samples_split=5, splitter='random')
```

```
In [27]: ▶ print(accuracy_score(y_train, dtc.predict(X_train)))
dtc_acc = accuracy_score(y_test, dtc.predict(X_test))
print(dtc_acc)
y_pred = dtc.predict(X_test)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0.967032967032967
```

```
0.9473684210526315
```

```
[[67  0]
```

```
 [ 6 41]]
```

	precision	recall	f1-score	support
0	0.92	1.00	0.96	67
1	1.00	0.87	0.93	47
accuracy			0.95	114
macro avg	0.96	0.94	0.94	114
weighted avg	0.95	0.95	0.95	114

## SVC

```
In [28]: ▶ from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
svc = SVC(probability=True)

parameters = {
    'gamma': [0.0001, 0.001, 0.01, 0.1],
    'C': [0.01, 0.05, 0.5, 0.1, 1, 10, 15, 20]
}
grid_search = GridSearchCV(svc, parameters)
grid_search.fit(X_train, y_train)
```

Out[28]:

```
▶ GridSearchCV
  ▶ estimator: SVC
    ▶ SVC
```

```
In [29]: ▶ grid_search.best_params_
```

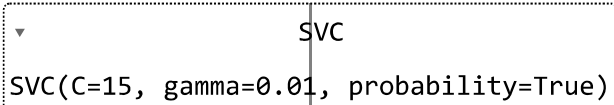
Out[29]: {'C': 15, 'gamma': 0.01}



In [30]: `grid_search.best_score_`

Out[30]: 0.9802197802197803

In [31]: `svc = SVC(C=15, gamma=0.01, probability=True)`  
`svc.fit(X_train, y_train)`

Out[31]: 

In [32]: `y_pred = svc.predict(X_test)`

In [33]: `print(accuracy_score(y_train, svc.predict(X_train)))`  
`svc_acc = accuracy_score(y_test, svc.predict(X_test))`  
`print(svc_acc)`  
`y_pred = svc.predict(X_test)`  
`print(confusion_matrix(y_test, y_pred))`  
`print(classification_report(y_test, y_pred))`

0.989010989010989

0.9824561403508771

[[67 0]

[ 2 45]]

	precision	recall	f1-score	support
0	0.97	1.00	0.99	67
1	1.00	0.96	0.98	47
accuracy			0.98	114
macro avg	0.99	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

## RandomForestClassifier

```
In [34]: from sklearn.ensemble import RandomForestClassifier

rand_clf = RandomForestClassifier(criterion = 'entropy', max_depth = 10, n_estimators=100)
rand_clf.fit(X_train, y_train)
```

```
Out[34]:
RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=10, max_features=0.5,
                        min_samples_leaf=2, min_samples_split=3,
                        n_estimators=130)
```

```
In [35]: y_pred = rand_clf.predict(X_test)
```

```
In [36]: print(accuracy_score(y_train, rand_clf.predict(X_train)))
rand_clf_acc = accuracy_score(y_test, rand_clf.predict(X_test))
print(rand_clf_acc)
y_pred = rand_clf.predict(X_test)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0.9978021978021978
```

```
0.9824561403508771
```

```
[[66  1]
 [ 1 46]]
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	67
1	0.98	0.98	0.98	47
accuracy			0.98	114
macro avg	0.98	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

## XGBClassifier

```
In [37]: pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\ishan\anaconda3\lib\site-packages (2.1.0)
```

```
Requirement already satisfied: numpy in c:\users\ishan\anaconda3\lib\site-packages (from xgboost) (1.23.5)
```

```
Requirement already satisfied: scipy in c:\users\ishan\anaconda3\lib\site-packages (from xgboost) (1.10.0)
```

```
Note: you may need to restart the kernel to use updated packages.
```

```
In [38]: ▶ from xgboost import XGBClassifier

xgb = XGBClassifier(objective = 'binary:logistic', learning_rate = 0.01, n

xgb.fit(X_train, y_train)
```

Out[38]:

```
XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=0.01, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
```

```
In [39]: ▶ print(accuracy_score(y_train, xgb.predict(X_train)))
xgb_acc = accuracy_score(y_test, xgb.predict(X_test))
print(xgb_acc)
y_pred = xgb.predict(X_test)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0.9934065934065934
```

```
0.956140350877193
```

```
[[65  2]
 [ 3 44]]
```

	precision	recall	f1-score	support
0	0.96	0.97	0.96	67
1	0.96	0.94	0.95	47
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

## Model Evaluation

```
In [40]: ▶ models = pd.DataFrame({
    'Model': ['Logistic Regression', 'Decision Tree Classifier', 'SVM', 'F'],
    'Score': [100*round(log_reg_acc,4), 100*round(dtc_acc,4), 100*round(sv
              100*round(xgb_acc,4)]
    })
models.sort_values(by = 'Score', ascending = False)
```

Out[40]:

	Model	Score
2	SVM	98.25
3	Random Forest Classifier	98.25
0	Logistic Regression	96.49
4	XgBoost	95.61
1	Decision Tree Classifier	94.74